

Design Specification document for Product Categorization

Version 1.0

1. Introduction

This document provides a comprehensive design for a machine learning system to predict product categories (e.g., "Sports & Outdoors," "Fashion") from its associated structured and unstructured product metadata.

1.1 Problem Statement

Objective: Given product attributes (title, description, features, manufacturer, SKU, details), classify each product into the correct category before building a proof-of-concept pipeline.

Challenges:

- Mixed data types (text, numerical, categorical, nested JSON).
- Key fields like 'Price,' 'Description,' or 'Features' may be missing.
- Some categories may dominate (in terms of number of occurrences) and can bias the model.
- Descriptions/titles contain unnecessary jargon, special characters, or inconsistent formatting.
- Need to derive structured features (e.g., weight, volume) from unstructured 'Details' while accounting for unit mismatches.
- Grouping and mapping metadata such as weight and volume is not straightforward due to inconsistent naming conventions.
- Need to accommodate added categories, if any, in the future.

1.2 Assumptions

- Features such as 'Description,' 'Details,' and 'Manufacturer' are mostly populated.
- The algorithm should prioritize accuracy over speed (can exceed ideal real-time end-to-end time).

- 'Price,' 'Weight,' 'Volume,' 'Material,' and 'Size' can be sparse but useful if imputed.
- The text fields and 'Manufacturer' field contain discriminative keywords (e.g., "Brake" → 'Automotive').
- The model should handle ~10k product categories with moderate resource usage and be scalable to ~1 million categories.
- Text fields contribute more significantly to predictions than numerical fields.

1.3 Open Questions

| Question | Possible Approaches |
|--|---|
| How to handle highly skewed numerical variables? | Log-transform & standardize, bin into quantiles, or drop as outliers. |
| Is NER worth the overhead? | Test spaCy NER using text based data |
| Best way to merge dimension fields? | Use regex to extract and standardize for units |
| How to handle similar categories like sports and fitness | Club them under a single category |
| Which model has ideal trade off between accuracy and memory requirements | Iterate through different models, starting from the base models and compare |

2. Functional Requirements

- Extracting and preprocessing necessary fields, including normalization and standardization.
- Identifying appropriate features for model training.
- Tokenization, lemmatization, and stop word removal.

- Grouping or generating features for training.
- Handling missing values, outliers, and data skewness.
- Addressing data imbalances.
- Hyperparameter tuning.
- Developing a classification model using ML algorithms.

3. Non-Functional Requirements

- Scalability considerations.
- Efficient memory management.
- Logging and tracking models and versions.
- Security and privacy.

4. Feature Engineering

| Feature Type | Examples | Encoding |
|-------------------|----------------------------|-----------------------------------|
| Numerical | Price, Weight, Volume | StandardScaler |
| Categorical | Manufacturer, SKU | One-Hot (if <50 cats), else Label |
| Text | Combined text | TF-IDF or BERT |
| Derived (via NER) | text ner, manufacturer NER | One-Hot |

4.1 Approaches

- **Pattern Matching:** Extract structured information from JSON using tools like regex.

- **NER (Named Entity Recognition):** Extract brand names and product materials from all text fields and the manufacturer field.
- **Dimensionality Reduction:** Use PCA (Principal Component Analysis) or ICA (Independent Component Analysis).
- **Feature Selection:** Use techniques like correlation matrix analysis, dimensionality reduction, and feature importance from base models.

5. Model Development

5.1 Class Imbalance Handling

- **Data-Level:** Use SMOTE for minority classes.
- **Algorithm-Level:** Utilize class weights in XGBoost and Random Forest.

5.2 Text Processing Methods

- **Basic Approaches:** TF-IDF or Bag-of-Words for lightweight solutions.
- **Advanced Approaches:** Use BERT embeddings for better contextual understanding.
- **NER:** Utilize spaCy's built-in models (e.g., `en_core_web_trf`) to derive structured features.

5.3 Baseline Models

- **Random Forest:** Handles heterogeneous features well.
 - Key hyperparameters: `n_estimators`, `max_depth`, `min_samples_split`, `min_samples_leaf`
- **XGBoost:** Optimized for imbalanced data.
 - Key hyperparameters: `n_estimators`, `max_depth`, `subsample`
- **SVM:** Suitable for small and complex datasets.
 - Key hyperparameters: `kernel`, `gamma`

5.4 Advanced Models

- **BERT:**

- Use `distilbert-base-uncased` for faster training.
- Full BERT model for improved accuracy.
- Key hyperparameters: `learning_rate`, `batch_size`, `embedding_size`
- **Neural Networks:**
 - LSTMs or RNNs as a fallback if BERT is too resource-intensive.
 - Key hyperparameters: `learning_rate`, `batch_size`

6. Evaluation

| Metric | Purpose |
|-----------------------------|--|
| Accuracy, Balanced Accuracy | Overall performance |
| F1-Score (per-class) | Better for imbalanced data. |
| Confusion Matrix | Identify misclassifications |
| Precision & Recall | Can prioritize one over the other based on the deployment considerations |
| Cross validation accuracy | Stratified K-Fold to ensure category distribution balance. |

7. Deployment & Monitoring

7.1 API

- **FastAPI Endpoint:**

- `POST /predict` (Input: JSON, Output: predicted category)

7.2 Containerization

- **Technologies:** Docker + Kubernetes for scalability.
- **Inference Optimization:** Convert models to ONNX or TensorRT for efficiency.

7.3 Maintenance & Monitoring

- **Retraining Strategy:** Evaluate continuously with added data and retrain as needed.
- **Logging:** Store predictions for continuous improvement.

8. Risks & Mitigation

| Risk | Severity | Possible Mitigations |
|-----------------------------|----------|---|
| Poor text feature quality | High | Augment with NER/regex features. |
| GPU out-of-memory with BERT | Medium | Use knowledge distillation techniques, decrease batch size and divide the dataset into chunks etc |
| Category drift over time | High | Regular retraining |
| Overfitting | High | Regularization, cross validation |
| Imbalanced datasets | Medium | Use upsample or downsample techniques |
| Order induced Bias | Medium | Randomize the data rows/samples |

9. Alternative/additional Approaches

9.1 Hyperparameter Tuning

- Use grid search or random search for optimization.

9.2 Feature Engineering Enhancements

- Incorporate additional features like material and size.
- Use FastText for handling out-of-vocabulary words if speed is a higher priority than performance.

9.3 Others

- Explore alternative techniques for imbalance handling
- Experiment with different embedding strategies beyond TF-IDF and BERT.
- Experiment with zero shot learning models
- Try a combination of models, if appropriate
- Check SHAP for better model interpretability

10. Deployment Strategies

- Consider serverless deployment using FastAPI + Docker for cost-effective scaling
- AWS Sagemaker for better control and ease of deployment

11. Implementation Plan

N/A

12. Tools & Libraries

- **Data Processing:** Pandas, spaCy, Regex.
- **Machine Learning:** Scikit-learn, XGBoost, Hugging Face, TensorFlow, torch, Transformers

- **Model Versioning & Tracking:** MLflow, Weights & Biases.
- **Deployment:** FastAPI, Docker.

This structured document provides a clear roadmap for developing and deploying a robust product categorization model.